

# Retail\_exploratory.pynb

March 9, 2018

Let's do some exploratory analysis on store sales data. Let's load and combine the data first.

```
In [8]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from datetime import datetime
from pandas import Series, DataFrame, concat

In [9]: def build_dataset():
    #LOAD DATA
    features_data = pd.read_csv('~/.Projects/Store_sales/Features_data_set.csv').fillna('VK')
    sales_data = pd.read_csv('~/.Projects/Store_sales/sales_data_set.csv').fillna('VK')
    stores_data = pd.read_csv('~/.Projects/Store_sales/stores_data_set.csv').fillna('VK')

    #MERGE DATA INTO A SINGLE DATA SET
    features_data.drop('IsHoliday', axis = 1, inplace = True)
    data = pd.merge(sales_data, stores_data, on='Store')
    data = pd.merge(data, features_data, on=['Store', 'Date'])
    del(features_data, sales_data, stores_data)

    #REPLACE DATE STRING WITH DATETIME
    data['Date'] = pd.Series([datetime.strptime(d, '%d/%m/%Y') for d in data['Date']])

    #ADD YEAR, MONTH AND WEEK COLUMNS
    data['Year'] = pd.Series([t.year for t in data['Date']])
    data['Month'] = pd.Series([t.month for t in data['Date']])
    '''there is a slight inconsistency here because of a leap year and the
    defnintion of the sales week depending on what week day the year starts;
    however given the dates are all fridays the datetime week function does
    a good job assigning week numbers'''
    data['Week'] = pd.Series([t.week for t in data['Date']])

    return data

In [3]: data = build_dataset()
```

Let's take a look now what this data set contains.

```
In [4]: data.describe(include = 'all')
```

```
Out[4]:
```

	Store	Dept	Date	Weekly_Sales	\
count	421570.000000	421570.000000	421570	421570.000000	
unique	NaN	NaN	143	NaN	
top	NaN	NaN	2011-12-23 00:00:00	NaN	
freq	NaN	NaN	3027	NaN	
first	NaN	NaN	2010-02-05 00:00:00	NaN	
last	NaN	NaN	2012-10-26 00:00:00	NaN	
mean	22.200546	44.260317	NaN	15981.258123	
std	12.785297	30.492054	NaN	22711.183519	
min	1.000000	1.000000	NaN	-4988.940000	
25%	11.000000	18.000000	NaN	2079.650000	
50%	22.000000	37.000000	NaN	7612.030000	
75%	33.000000	74.000000	NaN	20205.852500	
max	45.000000	99.000000	NaN	693099.360000	

	IsHoliday	Type	Size	Temperature	Fuel_Price	\
count	421570	421570	421570.000000	421570.000000	421570.000000	
unique	2	3	NaN	NaN	NaN	
top	False	A	NaN	NaN	NaN	
freq	391909	215478	NaN	NaN	NaN	
first	NaN	NaN	NaN	NaN	NaN	
last	NaN	NaN	NaN	NaN	NaN	
mean	NaN	NaN	136727.915739	60.090059	3.361027	
std	NaN	NaN	60980.583328	18.447931	0.458515	
min	NaN	NaN	34875.000000	-2.060000	2.472000	
25%	NaN	NaN	93638.000000	46.680000	2.933000	
50%	NaN	NaN	140167.000000	62.090000	3.452000	
75%	NaN	NaN	202505.000000	74.280000	3.738000	
max	NaN	NaN	219622.000000	100.140000	4.468000	

	MarkDown1	MarkDown2	MarkDown3	MarkDown4	MarkDown5	CPI	\
count	421570	421570	421570	421570	421570	421570.000000	
unique	2278	1500	1663	1945	2294	2145.000000	
top	VK	VK	VK	VK	VK	129.855533	
freq	270889	310322	284479	286603	270138	711.000000	
first	NaN	NaN	NaN	NaN	NaN	NaN	
last	NaN	NaN	NaN	NaN	NaN	NaN	
mean	NaN	NaN	NaN	NaN	NaN	NaN	
std	NaN	NaN	NaN	NaN	NaN	NaN	
min	NaN	NaN	NaN	NaN	NaN	NaN	
25%	NaN	NaN	NaN	NaN	NaN	NaN	
50%	NaN	NaN	NaN	NaN	NaN	NaN	
75%	NaN	NaN	NaN	NaN	NaN	NaN	
max	NaN	NaN	NaN	NaN	NaN	NaN	

Unemployment	Year	Month	Week
--------------	------	-------	------

count	421570.000	421570.000000	421570.000000	421570.000000
unique	349.000	NaN	NaN	NaN
top	8.099	NaN	NaN	NaN
freq	5152.000	NaN	NaN	NaN
first	NaN	NaN	NaN	NaN
last	NaN	NaN	NaN	NaN
mean	NaN	2010.968591	6.449510	25.826762
std	NaN	0.796876	3.243217	14.151887
min	NaN	2010.000000	1.000000	1.000000
25%	NaN	2010.000000	4.000000	14.000000
50%	NaN	2011.000000	6.000000	26.000000
75%	NaN	2012.000000	9.000000	38.000000
max	NaN	2012.000000	12.000000	52.000000

In [5]: data.dtypes

```
Out[5]: Store          int64
Dept          int64
Date          datetime64[ns]
Weekly_Sales  float64
IsHoliday     bool
Type          object
Size          int64
Temperature   float64
Fuel_Price    float64
MarkDown1     object
MarkDown2     object
MarkDown3     object
MarkDown4     object
MarkDown5     object
CPI           object
Unemployment   object
Year          int64
Month         int64
Week          int64
dtype: object
```

Now it would be interesting to take a look at an individual store just to make sense we are clear as to what the building block of the data set represents. Let's select store 1 and see what we got. We will show two separate graphs just to keep the spatial presentation under control:

```
In [11]: select_store = 1
         select_dept = 1
         data_df = data.loc[(data['Store'] == select_store) & (data['Dept'] == select_dept)]

         #PLOT SELECTED COLUMNS FOR THE SELECTED STORE - FIRST GROUP
         values = data_df.values
         to_plot = [ 3, 4, 5, 6, 7, 8, 9]
         i = 1
```

```

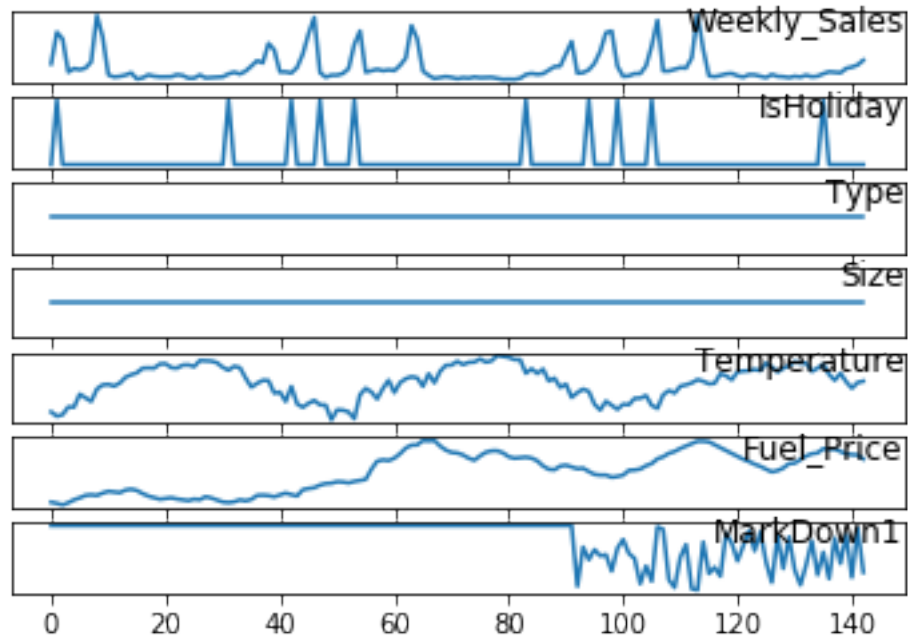
plt.figure()
for item in to_plot:
    plt.subplot(len(to_plot), 1, i)
    plt.plot(values[:, item])
    plt.yticks([])
    plt.title(data_df.columns[item], y=0.5, loc='right')
    i += 1
plt.show()

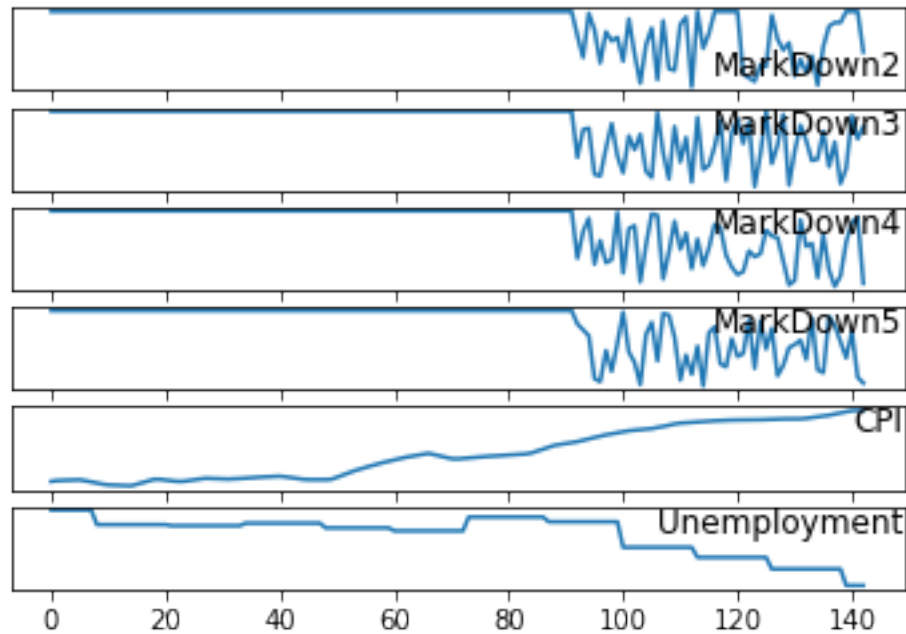
```

```

#AND THEN SECOND GROUP
to_plot = [10, 11, 12, 13, 14, 15]
i = 1
# plot each column
plt.figure()
for item in to_plot:
    plt.subplot(len(to_plot), 1, i)
    plt.plot(values[:, item])
    plt.yticks([])
    plt.title(data.columns[item], y=0.5, loc='right')
    i += 1
plt.show()

```





I think it gives some good idea as to what should be useful here - weekly sales for sure, IsHoliday factor perhaps, and maybe some other values for hygiene (e.g. CPI, unemployment, etc.) but perhaps these won't really swing the results. I don't quite understand the nature of Markdowns and what they represent and I don't think I will use them in the analysis anyway for this reason, and because they are only only available for 1/3 of the time series.

I also think that given the seasonal nature of the data, it would be important to add number of the week to the data set. Sales spikes occur in week 47 and 51 and we should be able to capture this.

Now let's look at some other useful things.

```
In [19]: #THE NUMBER OF UNIQUE STORES
print('The number of unique stores:')
data['Store'].nunique()
```

The number of unique stores:

Out[19]: 45

```
In [20]: #PLOT STORES VS NUMBER OF DEPARTMENTS FOR EACH STORE
def store_plot():

    x = data['Store'].unique()
    y = []
    for i in range(len(x)):
        temp_df = data.loc[data['Store'] == i+1]
        y.append(temp_df['Dept'].nunique())
```

```
i += 1
```

```
plt.figure()
plt.scatter(x, y, alpha=0.5)
plt.xlabel('stores')
plt.ylabel('departments')
plt.show()
```

```
In [21]: #ANNUAL SALES PER STORE
```

```
df_2010 = data[data['Date'].isin(pd.date_range("2010-01-01", "2010-12-31"))]
df_2011 = data[data['Date'].isin(pd.date_range("2011-01-01", "2011-12-31"))]
df_2012 = data[data['Date'].isin(pd.date_range("2012-01-01", "2012-12-31"))]


x = data['Week'].unique()
x = df_2010['Week'].unique()
y = df_2010.groupby('Week')['Weekly_Sales'].sum()
```

```
In [22]: #build table aggregating sales by year and week
```

```
annual_sales_df = pd.concat([df_2010.groupby('Week')['Weekly_Sales'].sum(), df_2011.groupby('Week')['Weekly_Sales'].sum(), df_2012.groupby('Week')['Weekly_Sales'].sum()])
annual_sales_df = pd.concat([annual_sales_df, df_2012.groupby('Week')['Weekly_Sales'].sum()])
annual_sales_df.columns = ['2010', '2011', '2012']
```

And now let's take a look at the plot showing all aggregated sales by week of the year.


```
In [24]: plt.figure()
plt.plot(annual_sales_df['2010'])
plt.plot(annual_sales_df['2011'])
plt.plot(annual_sales_df['2012'])
plt.legend(loc='upper left')
plt.show()
```



It would be helpful at this stage to also look separately at each year aggregate sales, 2010, 2011 and 2012:

```
In [25]: df_2010.groupby(['Week', 'Store']).sum()['Weekly_Sales'].unstack().plot(legend = None,
plt.title('Weekly sales by store in 2010')
plt.xlim(1,52)
```

```
Out[25]: (1, 52)
```



```
In [26]: df_2011.groupby(['Week', 'Store']).sum()['Weekly_Sales'].unstack().plot(legend = None)
plt.title('Weekly sales by store in 2011')
plt.xlim(1,52)
```

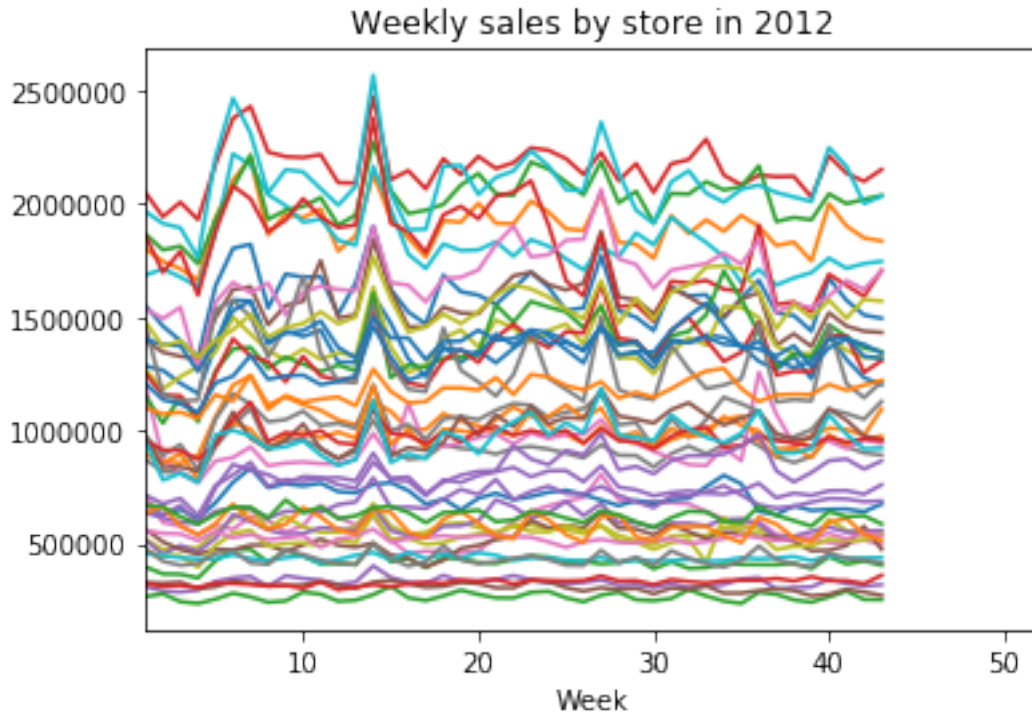
```
Out[26]: (1, 52)
```





```
In [27]: df_2012.groupby(['Week', 'Store']).sum()['Weekly_Sales'].unstack().plot(legend = None)
plt.title('Weekly sales by store in 2012')
plt.xlim(1,52)
```

```
Out[27]: (1, 52)
```



I think these graphs give a pretty complete picture of the data: we have incomplete year 2010 (a few weeks in the beginning) and 2012 (~7 weeks at the end, including all of the holiday season). Now it would be interesting to see if a model can be built that predicts the usual seasonal fluctuation based on the observations from 2010 and 2011.

I am not going to bother with building a train/test split this time, but instead will use the full data set as for training. It's pretty clear what we can expect to see in that gap at the end of 2012 graph.

The purpose here is only to set up a model that would correctly infer the typical sales patterns surrounding the holiday season. LSTM is an obvious model candidate. Given LSTM choice, it would be interesting to:

- make a series of predictions in one go, i.e. make one prediction of several weeks, as opposed to a week by week prediction
- experiment with adding different independent variables and their effect on the model performance

This is exactly what I will try next.